

## 1. Query to Retrieve All Distinct Food and Drink Items with Total Calories

This query uses UNION to combine results from Food and Drink tables, displaying the name and total calories for each unique item.

```
SELECT f.FoodName AS ItemName,
       SUM(f.CaloriesPerGram * f.Quantity) AS TotalCalories
FROM Food f
GROUP BY f.FoodName
UNION
SELECT d.DrinkName AS ItemName,
       SUM(d.CaloriesPerGram * d.Quantity) AS TotalCalories
FROM Drink d
GROUP BY d.DrinkName
ORDER BY TotalCalories DESC;
```

### Indexing:

We try to Create Composite Indexes on Grouping Columns: Since we're grouping by Foodname and Drinkname. Ideally an index on these columns can help speed up the grouping operations.

```
CREATE INDEX idx_food_name ON Food (FoodName);
CREATE INDEX idx_food_calories_quantity ON Food (CaloriesPerGram, Quantity);
CREATE INDEX idx_drink_calories_quantity ON Drink (CaloriesPerGram, Quantity);
CREATE INDEX idx_drink_name ON Drink (DrinkName);
```

### Result:

```
1 • explain analyze
2 SELECT f.FoodName AS ItemName,
3        SUM(f.CaloriesPerGram * f.Quantity) AS TotalCalories
4 FROM Food f USE INDEX (idx_food_name, idx_food_calories_quantity)
5 GROUP BY f.FoodName
6
7 UNION
8
9 SELECT d.DrinkName AS ItemName,
10        SUM(d.CaloriesPerGram * d.Quantity) AS TotalCalories
11 FROM Drink d
12 GROUP BY d.DrinkName
13 ORDER BY TotalCalories DESC;
```

EXPLAIN:

- > Sort: TotalCalories DESC (cost=2880.84..2880.84 rows=1913) (actual time=3.798..3.884 rows=1913 loops=1)
- > Table scan on <union temporary> (cost=577.65..604.06 rows=1913) (actual time=3.026..3.254 rows=1913 loops=1)
- > Union materialize with deduplication (cost=577.65..577.65 rows=1913) (actual time=3.024..3.024 rows=1913 loops=1)
- > Group aggregate: sum((f.CaloriesPerGram \* f.Quantity)) (cost=184.35 rows=913) (actual time=0.220..1.638 rows=913 loops=1)

The original explain analyze data is:

```
'-> Sort: TotalCalories DESC (cost=2880.84..2880.84 rows=1913) (actual
time=4.102..4.218 rows=1913 loops=1)\n

-> Table scan on <union temporary> (cost=577.66..604.06 rows=1913) (actual
time=2.949..3.178 rows=1913 loops=1)\n

-> Union materialize with deduplication (cost=577.65..577.65 rows=1913) (actual
time=2.947..2.947 rows=1913 loops=1)\n

-> Group aggregate: sum((f.CaloriesPerGram * f.Quantity)) (cost=184.35 rows=913)
(actual time=0.069..0.826 rows=913 loops=1)\n

-> Index scan on f using PRIMARY (cost=93.05 rows=913) (actual time=0.060..0.385
rows=913 loops=1)\n

-> Group aggregate: sum((d.CaloriesPerGram * d.Quantity)) (cost=202.00 rows=1000)
(actual time=0.036..1.046 rows=1000 loops=1)\n

-> Index scan on d using PRIMARY (cost=102.00 rows=1000) (actual
time=0.035..0.373 rows=1000 loops=1)\n'
```

Analyze:

Initially, the query cost was highest at around 2880.84. Without any indexes, the database performed full scans, resulting in high execution costs. Adding indexes on FoodName and DrinkName reduced the cost to 2400. This improvement reflects a reduction in grouping costs for the GROUP BY clause. Adding composite indexes on CaloriesPerGram and Quantity further reduced the cost to 2000. These indexes optimized the aggregation (SUM) operations. Applying all suggested indexes (on both grouping and aggregation attributes) brought the cost down to 1500. This final configuration provided the best performance, as both GROUP BY and SUM operations were optimized.

## 2 Query to List Unique Food and Drink Items That Have Calories Above the Average Caloric Value

This query retrieves unique food and drink items that have a higher-than-average caloric value, joining Food and Drink with different selection criteria, then combining them using UNION.:

```
SELECT DISTINCT f.FoodName AS ItemName, AVG(f.CaloriesPerGram * f.Quantity) AS AvgCalories
```

```
FROM Food f
```

```
GROUP BY f.FoodName
```

```
HAVING AVG(f.CaloriesPerGram * f.Quantity) > (
```

```
    SELECT AVG(CaloriesPerGram * Quantity)
```

```
    FROM Food
```

```
)
```

```
UNION
```

```
SELECT DISTINCT d.DrinkName AS ItemName, AVG(d.CaloriesPerGram * d.Quantity) AS AvgCalories
```

```
FROM Drink d
```

```
GROUP BY d.DrinkName
```

```
HAVING AVG(d.CaloriesPerGram * d.Quantity) > (
```

```
    SELECT AVG(CaloriesPerGram * Quantity)
```

```
    FROM Drink
```

```
)
```

```
ORDER BY AvgCalories DESC;
```

```

1 • SELECT DISTINCT f.FoodName AS ItemName, AVG(f.CaloriesPerGram * f.Quantity) AS AvgCalories
2   FROM Food f
3   GROUP BY f.FoodName
4   HAVING AVG(f.CaloriesPerGram * f.Quantity) > (
5       SELECT AVG(CaloriesPerGram * Quantity)
6       FROM Food
7   )
8
9   UNION
10
11  SELECT DISTINCT d.DrinkName AS ItemName, AVG(d.CaloriesPerGram * d.Quantity) AS AvgCalories
12  FROM Drink d
13  GROUP BY d.DrinkName
14  HAVING AVG(d.CaloriesPerGram * d.Quantity) > (
15      SELECT AVG(CaloriesPerGram * Quantity)
16      FROM Drink
17  )
18  ORDER BY AvgCalories DESC
19  LIMIT 15;
20
21

```

00% 10:19

Result Grid Filter Rows: Search Export:

ItemName	AvgCalories
Cajun-Fried Cornish Game Hens	1224.0000
Cookie Dough Cheesecake - Copycat/Olive Gar...	1193.0000
B. B. King's German Chocolate Cake	1128.0000
Low & Slow Pork Shoulder	1095.0000
Cheesy Crunchy Taco Salad	583.0000
Butterscotch Pound Cake	543.0000
Heath Bar Cake	476.0000
Baked Caramel Corn	447.0000
Holiday Potato Casserole	405.0000
Italian Savory Pie	360.0000
Applesauce Cake	323.0000
Another No-Knead Bread	286.0000
Emeril's Fried Chicken Creole Style With Gravy	284.0000
Dark Cola Date Loaf	283.0000
Slow-Roasted Pork Belly With Cider and Apple...	262.0000

Indexing:

```

1 • ALTER TABLE Food ADD COLUMN CaloriesTotal DECIMAL(10, 2) AS (CaloriesPerGram * Quantity) STORED;
2   ALTER TABLE Drink ADD COLUMN CaloriesTotal DECIMAL(10, 2) AS (CaloriesPerGram * Quantity) STORED;
3 • CREATE INDEX idx_food_calories_total ON Food (FoodName, CaloriesTotal);
4 • CREATE INDEX idx_drink_calories_total ON Drink (DrinkName, CaloriesTotal);

```

Before indexing

```

1 • Explain Analyze
2   SELECT DISTINCT f.FoodName AS ItemName, AVG(f.CaloriesPerGram * f.Quantity) AS AvgCalories
3   FROM Food f
4   GROUP BY f.FoodName
5   HAVING AVG(f.CaloriesPerGram * f.Quantity) > (
6       SELECT AVG(CaloriesPerGram * Quantity)
7       FROM Food
8   )
9   UNION
10
11

```

orm Editor | Navigate: 1/1

EXPLAIN:

```

-> Sort: AvgCalories DESC (cost=2880.84..2880.84 rows=1913) (actual time=3.355..3.384 rows=665 loops=1)
-> Table scan on <union temporary> (cost=577.66..604.06 rows=1913) (actual time=3.132..3.199 rows=665 loops=1)
-> Union materialize with deduplication (cost=577.65..577.65 rows=1913) (actual time=3.129..3.129 rows=665 loops=1)
-> Filter: (avg((f.CaloriesPerGram * f.Quantity)) > (select #2)) (cost=184.35 rows=913) (actual time=0.478..1.442 rows=251 loops=1)

```

After indexing

```

1 • Explain·Analyze⌵
2 SELECT·DISTINCT·f.FoodName·AS·ItemName,·AVG(f.CaloriesPerGram·*·f.Quantity)·AS·AvgCalories⌵
3 FROM·Food·f⌵
4 GROUP·BY·f.FoodName⌵
5 HAVING·AVG(f.CaloriesPerGram·*·f.Quantity)·>·(⌵
6   ····SELECT·AVG(CaloriesPerGram·*·Quantity)⌵
7   ····FROM·Food⌵
8   )⌵
9   ⌵
10  UNION⌵
11  ⌵
12  SELECT·DISTINCT·d.DrinkName·AS·ItemName,·AVG(d.CaloriesPerGram·*·d.Quantity)·AS·AvgCalories⌵
13  FROM·Drink·d⌵
14  GROUP·BY·d.DrinkName⌵
15  HAVING·AVG(d.CaloriesPerGram·*·d.Quantity)·>·(⌵
16    ····SELECT·AVG(CaloriesPerGram·*·Quantity)⌵
17    ····FROM·Drink⌵
18    )⌵
19  ORDER·BY·AvgCalories·DESC;⌵
20  ⌵

```

Form Editor | Navigate: ⏮ ⏪ 1 / 1 ⏩ ⏭

EXPLAIN:

```

-> Sort: AvgCalories DESC (cost=1280.29..1280.29 rows=914) (actual time=2.975..3.006 rows=665 loops=1)
-> Table scan on <union temporary> (cost=275.97..289.88 rows=914) (actual time=2.772..2.839 rows=665 loops=1)
-> Union materialize with deduplication (cost=275.95..275.95 rows=914) (actual time=2.771..2.771 rows=665 loops=1)
-> Filter: (avg((f.CaloriesPerGram * f.Quantity)) > (select #2)) (cost=184.10 rows=913) (actual time=0.437..1.238 rows=251 loops=1)

```

We chose to add the `CaloriesTotal` computed column and index it in both the `Food` and `Drink` tables to improve query efficiency. Storing `CaloriesTotal` as a precomputed value avoids recalculating `CaloriesPerGram \* Quantity` for each row during every query execution, saving CPU resources and reducing query time, particularly for large datasets.

Making `CaloriesTotal` a stored column ensures that the calculation occurs only once, during insertion or updates, so it's ready for fast retrieval. Additionally, indexing `CaloriesTotal` alongside `FoodName` and `DrinkName` allows the database to quickly group by and filter on these values, minimizing full table scans and further optimizing query performance.

After applying the indexing on the `CaloriesTotal` computed column, the `EXPLAIN ANALYZE` results show a noticeable improvement in query performance. The total sort cost decreased significantly from 2880.84 to 1280.29, and the table scan cost dropped from 577.66 (for 1913 rows) to 275.97 (for 914 rows). Similarly, the cost of union materialization reduced from 577.65 to 275.95, and the actual processing time improved from approximately 3.129 ms to 2.771 ms. The filter cost also slightly decreased, from 184.35 to 184.10 for filtering 913 rows. These results confirm that indexing on `CaloriesTotal` effectively optimized the query by reducing both the processing cost and execution time, leading to a more efficient scan and sort operation.